

Election Algorithm in Distributed Systems

Bully Algorithm



Bully Algorithm Main Characteristics

- **Operating Assumptions**

- Synchronous system
 - All messages arrive within T_M units transmission of time.
 - A reply is dispatched within P_p units of processing time after the receipt of a message.
 - If no response is received in $2 \times T_M + P_p$, the node is assumed to be faulty
 - Node crashed
 - Attribute = Process ID
 - Each process knows all the other processes in the system
 - Therefore, processes know each others' IDs
-



The Bully Algorithm

When any process, P, notices that the coordinator is no longer responding it initiates an election:


1. P sends an *election* message to all processes with higher id numbers.
 2. If no one responds, P wins the election and becomes coordinator.
 3. If a higher process responds, it takes over.
 - Process P's job is done.
-



The Bully Algorithm



- At any moment, a process can receive an **election** message from one of its lower-numbered colleagues.
 - The receiver sends an OK back to the sender and conducts its own election.
 - Eventually only the **bully process** remains.
 - Bully process becomes the new coordinator
 - The bully announces victory to all processes in the distributed group.
-

Bully Algorithm Example



- Process 4 notices 7 down.
- Process 4 holds an election.
- Process 5 and 6 respond, telling 4 to stop.
- Now 5 and 6 each hold an election.

Bully Algorithm Example



- Process 6 tells process 5 to stop
- Process 6 (the bully) wins and tells everyone
- If process 7 recovers, it restarts election process



Performance of Bully Algorithm

- **Best case scenario:** The process with the second highest id notices the failure of the coordinator and elects itself.
 - $N-2$ coordinator messages are sent.
 - Turnaround time is one message transmission time.
 - **Worst case scenario:** When the process with the least id detects the failure.
 - $N-1$ processes altogether begin elections, each sending messages to processes with higher ids.
 - The message overhead is $O(N^2)$.
 - Turnaround time is approximately 5 message transmission times if there are no failures during the run: election, answer, election, answer, coordinator
-

Ring Algorithm



Ring Algorithm – Basic Operation

- RA assumes that the processes are logically ordered in a ring **{implies a successor pointer and an active process list}** that is unidirectional.
- When any process, P, notices that the coordinator is no longer responding it initiates an election:
 1. P sends message containing **P's process ID** to the next available successor.



Ring Algorithm – Basic Operation

2. At each active process, the receiving process adds its process number to the list of processes in the message and forwards it to its successor.
 - Eventually, the message gets back to the sender.
3. The initial sender sends out a second message letting everyone know who the coordinator is **{the process with the highest number}** and indicating the current members of the active list of processes.